



**Hochschule
Zittau/Görlitz**
UNIVERSITY OF APPLIED SCIENCES

Master's Thesis

Something with XAI stuff and making it actually work

Master's Program in Computer Science

Jonathan Zopf

First Examiner: Prof. Dr. Jörg Lässig

Second Examiner: M. Sc. Falko Gawantka

Faculty of Electrical Engineering and Computer Science
University of Applied Sciences Zittau-Görlitz

Görlitz, July 2, 2026

Declaration of Originality

I hereby declare that I have written this thesis independently and have not used any sources other than those indicated. All passages that are taken verbatim or in spirit from published or unpublished sources are clearly marked as such. This work has not been submitted to any other examination board in the same or similar form.

Görlitz, July 2, 2026

A handwritten signature in black ink, appearing to read "Jonathan Zepf". The signature is stylized with a large, sweeping "Z" and a final flourish.

Contents

1	Introduction	1
2	Background	2
2.1	Introduction to LLM Explainability	2
2.1.1	Definition and Purpose	2
2.1.2	Transparency as Regulatory Requirement	2
2.1.3	Audience Groups	3
2.1.4	Aspects and Dimensions	3
2.1.5	LLM Explainability Challenges	4
2.2	Faithfulness of LLM Explanations	4
2.2.1	Definition and Significance	4
2.2.2	Estimation Approaches	5
2.3	Causal Concept Faithfulness Metric	8
2.3.1	Introduction	8
2.3.2	Mathematical Aspects	9
2.3.3	Causal Concept Faithfulness	11
2.3.4	Evaluation Pipeline	11
2.3.5	Estimating Causal Concept Faithfulness	13
2.3.6	The Case for this Metric	13
2.4	The Correlated Concepts Problem	13

Chapter 1

Introduction

sdasd

Chapter 2

Background

2.1 Introduction to LLM Explainability

2.1.1 Definition and Purpose

Large language models (LLMs) have seen increasing adoption in recent years for tasks such as question answering [3], writing [18], and software development [42], but also in high-stakes domains such as healthcare [23, 16, 13]. However, since LLMs are complex deep learning models with millions of neurons and billions of parameters, they are often considered black boxes [30]. This makes it challenging to comprehend and evaluate their decisions.

Explainable AI (XAI) is a research field that seeks to make artificial intelligence algorithms transparent and understandable [12], thereby making the inherent black box less opaque. Central to this research field is the concept of explainability, whose precise definition remains disputed [33]. The term is described in the relevant ISO standard as the “level of understanding of how the AI-based system came up with a given result” [14]. Explainability aims to justify decisions and demonstrate that they are correct and unbiased in order to build trust with users [2].

The term *interpretability* is often used in a similar context, although there is a subtle distinction between the two concepts. Interpretability concerns understanding the inner workings of an AI model, whereas explainability focuses on making the decision process retraceable [8, 2].

2.1.2 Transparency as Regulatory Requirement

The recent increase in the importance of XAI research has also been driven by new regulatory measures. While multiple countries have introduced their own legislative frameworks to govern the use of AI, such as China [45] and the United States [5], the European Union’s AI Act [31] can be regarded as particularly relevant to explainability. Under this regulation, individuals affected by high-risk AI decisions have the right to obtain clear and meaningful explanations regarding the AI system’s role and the main contributing factors to the decision. Furthermore, high-risk AI systems must provide transparency and allow for human oversight [31]. The *Right to Explanation*

derived from the General Data Protection Regulation (GDPR) is also relevant to AI systems within the European Union [27].

2.1.3 Audience Groups

Providing explanations for reasoning processes or decisions made by an AI system can increase trust in the system among end users and experts [40, 9], while also assisting researchers and developers in debugging and fine-tuning the model [37]. Moreover, as argued in Chapter 2.1.2, explainability can be required by regulations and auditors, thereby forming another stakeholder group for XAI systems [9, 37].

It has also been noted that different audiences perceive and consume explanations differently [10, 9]. For instance, developers may expect technical information about the inner workings of the system, whereas end users may benefit more from concise and understandable explanations in natural language. Adapting explanations generated by an AI system to the expectations of the recipient has therefore become a subject of recent research [11, 25].

2.1.4 Aspects and Dimensions

Herrera [9] points out that, whereas traditional AI models usually follow fixed deterministic rules when generating explanations, LLMs complicate this process due to their more complex and advanced capabilities. The author therefore proposes four dimensions of explainability for large language models:

- **Faithfulness:** Faithfulness describes whether the explanation corresponds to the internal reasoning of the model [15].
- **Truthfulness:** Truthfulness describes whether the explanation conforms to real-world facts derived from an external authoritative reference. A truthful explanation does not contain factual errors, omissions, or distortions [9].
- **Plausibility:** Plausibility describes the degree to which an explanation is convincing and appears reasonable to humans [44, 15].
- **Contrastivity:** Contrastivity measures how well the explanation highlights differences between competing pieces of information. In classification tasks, an XAI system with high contrastivity clearly outlines *why* it selected one class over another [9, 26, 35].

It should be noted that there is still no broad scientific consensus within the XAI community regarding which aspects are essential for an explanation to be considered good, since descriptions, terminology, and definitions of these factors vary considerably [29, 2, 19, 26]. This issue is further amplified by regulatory and standardization frameworks, which remain comparatively vague in this regard [31, 14]. Nevertheless, Herrera’s [9] desiderata of faithfulness and plausibility in particular are supported by other literature, including [15, 44, 38]. Other authors, such as Lopes et al. [19], refer to faithfulness as *fidelity* and divide plausibility into several human-centered sub-aspects.

2.1.5 LLM Explainability Challenges

According to Zhao et al. [43, 21], explainability research for LLMs is difficult because understanding their internal processes becomes infeasible due to their enormous size and complex structure. Traditional methods such as gradient-based attribution [36] and SHAP values [20] are impractical for state-of-the-art LLMs with billions of parameters. This complexity also hinders in-depth analysis, debugging, and diagnosis of these models. Furthermore, the opaque black-box nature of LLMs amplifies this problem [30]. Many mainstream commercial models are only accessible via APIs [24], meaning that their model weights are not publicly available to users or researchers.

Despite these limitations, several technical approaches have been developed to gain insight into the black box and retrace important reasoning pathways within an LLM. These approaches cannot be discussed in detail within the scope of this thesis, but overviews of the most prominent methods are provided in [43, 21].

At the same time, since LLMs are specifically designed for natural language processing tasks, they also provide an opportunity to generate explanations directly in natural human language [43]. Most notably, this includes *ad-hoc* explanations where the model generates a justification for its decision after giving the prediction and *Chain-of-Thought* (CoT) explanations, in which LLMs reconstruct their own reasoning process. Prompting an LLM to generate CoT explanations has furthermore been shown to improve performance on complex reasoning tasks [41].

2.2 Faithfulness of LLM Explanations

2.2.1 Definition and Significance

Faithfulness describes the degree to which the explanation generated by an AI-powered system is grounded in the model’s actual internal reasoning process [9, 15, 32]. A faithful explanation should therefore reflect the model’s true decision path rather than merely providing a plausible post-hoc rationalization [9].

LLMs have the ability to give those explanations in natural language [41]. But unlike other explainability methods, the faithfulness of these are not guaranteed [9, 15]. In particular, an LLM may generate a natural language explanation that sounds convincing to users (high plausibility) while being entirely fabricated and not grounded in the model’s actual reasoning process (low faithfulness). Such scenarios can pose significant risks and lead to misplaced trust in these systems [1, 22, 9].

In particular, research by Turpin et al. [38] demonstrated that LLMs can systematically conceal internalized biases while providing plausible yet unfaithful justifications. When deliberately influenced to rely on prejudicial reasoning in their experiments, the models often used social biases such as gender or race as primary decision factors, but rarely mentioned them explicitly in their explanations. Instead, they resorted to alternative justifications based on conduct or circumstances. This behavior is especially concerning in applications such as recidivism prediction, where fairness may be undermined [7, 15]. Consequently, it is crucial not only to verify

whether an explanation appears reasonable, but also whether it accurately reflects the actual reasoning process of the model, particularly in high-stakes domains.

2.2.2 Estimation Approaches

General Challenges

As discussed in Section 2.1.5, the complex and opaque nature of large language models poses unique challenges for explainability research. These challenges are particularly pronounced when assessing the faithfulness of CoT explanations. Whereas plausibility can be evaluated by human judges [9], faithfulness is inherently tied to the model’s internal reasoning process [28, 15].

Faithfulness vs Self-Consistency

Some have argued such as [28] that faithfulness can not be measured and only self-consistency can be. Present here the case, why we still talk about measuring or estimating faithfulness.

Early Faithfulness Work

One of the earliest works addressing the estimation of faithfulness in LLMs was conducted by Turpin et al. [38]. In addition to demonstrating that CoT explanations can systematically omit contributing decision factors, the authors proposed a quantification approach for unfaithfulness on the BBH dataset. Their method measures the drop in model accuracy when biasing features are added to inputs that the model fails to mention in its explanations. However, the approach is labor-intensive, as CoT explanations generated by the LLM must be manually inspected by human annotators to determine whether the biased feature is mentioned. According to the authors, the method only evaluates faithfulness with respect to relatively small input modifications, whereas a practically useful metric should also be able to investigate reasoning behavior over larger input changes.

Another early contribution was presented by Chen et al. [6], who introduced the concept of *counterfactual simulatability*. This concept describes the degree to which humans can infer “the model’s outputs on diverse counterfactuals of the explained input.” In other words, after reading an explanation to an initial question, a human should be able to predict how the model will answer closely related counterfactual yes-or-no questions. Based on this concept, the authors propose two metrics: *simulation generality*, which measures the diversity of counterfactuals relevant to the explanation, and *simulation precision*, which measures the extent to which human expectations align with the model’s actual outputs on those counterfactuals.

Their evaluation pipeline first selects a question from a dataset and prompts the LLM to generate both an answer and an explanation. This explanation is subsequently used to generate corresponding counterfactual questions, a task performed by another LLM in order to scale to larger datasets. Human evaluators then construct a mental model based on the original explanation and state their expectations

regarding the answers to the counterfactual questions. Finally, the counterfactuals are presented to the LLM and the generality and precision scores are computed. This approach is notable because it supports both ad-hoc explanations and CoT explanations. It also represents an early example of using auxiliary LLMs to generate counterfactuals, thereby inheriting the limitations associated with LLM-as-a-Judge approaches [24].

Counterfactual Edits

Another influential use of AI-generated counterfactuals was introduced by Atanasova et al. [4]. The authors laid the foundation for a line of work commonly referred to as *counterfactual edits* [28, 24], which has subsequently been adopted and extended by later metrics [34, 24]. In these approaches, the LLM’s response to a counterfactual input serves as the basis for estimating faithfulness.

For their method, Atanasova et al. trained an auxiliary model to insert words into the original input, thereby generating counterfactuals. Both the original input and the modified version are then provided to the evaluated LLM. The resulting predictions and the explanation for the counterfactual input are subsequently analyzed. If the prediction changes due to the inserted words and those words are not mentioned in the explanation, the explanation is considered unfaithful. A major advantage of this approach is that it does not require human annotators [28].

Nevertheless, several limitations remain. The inserted modifications may shift the model’s attention to unrelated aspects of the input, meaning that the prediction change is not necessarily caused by the inserted feature itself [4]. Siegel et al. [34] additionally identify two practical shortcomings of this implementation of counterfactual edits. First, the method does not investigate whether highly influential features are more likely to be mentioned than weakly influential ones. In the extreme case, an AI system could theoretically achieve perfect faithfulness simply by repeating the entire input text. Second, the binary impact measurement only evaluates whether the prediction probability crosses the 50% threshold, thereby ignoring potentially meaningful probability shifts that remain below this boundary.

Correlation-Based Metrics

To address these limitations, Siegel et al. [34] proposed the Correlational Counterfactual Test (CCT), which significantly extends the approach of Atanasova et al. [4]. At its core, CCT measures faithfulness as the correlation between intervention impact scores and explanation mention scores [24].

First, the LLM receives an unmodified question from a dataset. The input is then perturbed and presented to the LLM again, producing prediction probability distributions for both the original and counterfactual versions. The counterfactual generation procedure largely follows the earlier work by Atanasova et al., although an auxiliary LLM is additionally used to filter nonsensical interventions. The Total Variation Distance (TVD) between the original and perturbed prediction distributions is then calculated in order to quantify the impact of the intervention. Subsequently, the

generated explanation is analyzed using substring matching to determine whether the intervention is explicitly mentioned, resulting in a binary explanation score. This procedure is repeated across multiple dataset examples, after which the Pearson correlation coefficient between intervention impacts and explanation mentions is computed to obtain the final CCT score.

Since the metric builds directly upon the approach by Atanasova et al. [4], several limitations are inherited, particularly with respect to counterfactual generation. For example, the counterfactuals are restricted to the insertion of individual adjectives or adverbs, limiting their semantic diversity. Furthermore, the metric does not account for cases in which an explanation refers to a synonym or semantically equivalent concept rather than the exact inserted word.

Output Contribution Metrics

Parcalabescu and Frank [28] argue that many existing metrics measure self-consistency of outputs rather than genuine faithfulness. They therefore propose their own self-consistency metric, CC-SHAP, which does not depend on modifying the input. Their approach compares the degree to which input tokens contribute to both the generated answer and the explanation. By measuring the convergence between these contributions using Shapley values, the method produces a continuous score.

This approach offers several advantages. It avoids the need for input perturbations, does not require semantic evaluation of explanations, supports both CoT and ad-hoc explanations, reveals patterns of consistency and inconsistency at the token level, and does not require auxiliary AI models during evaluation. However, these benefits come at the cost of considerable computational overhead, which may exceed that of other approaches [28].

CoT Corruption Methods

Lanham et al. [17] investigated faithfulness by deliberately corrupting CoT reasoning in order to test hypotheses regarding CoT behavior. First, the model is prompted to reason step-by-step about a multiple-choice problem. Given the original CoT, the LLM then produces a final answer. The CoT is subsequently modified through several interventions, including truncation, the insertion of mistakes using an auxiliary model, paraphrasing of earlier reasoning steps followed by regeneration of the remainder, and the insertion of filler tokens. The model is then asked again to provide an answer based on the modified CoT.

If the prediction remains unchanged despite the intervention, the explanation is considered unfaithful. Since the metric primarily serves as a tool for evaluating research hypotheses about CoT reasoning, it remains unclear how directly comparable it is to other faithfulness metrics. Parcalabescu and Frank [28] further argue that this approach assumes that the model fundamentally relies on the CoT in order to produce the correct prediction. However, experiments conducted by Lanham et al. [17] themselves suggest that CoT explanations only marginally improve prediction accuracy compared to models without CoT reasoning [28], thereby raising questions

regarding the practical usefulness of this method.

Faithfulness Through Unlearning

A final approach is presented by Tutek et al. [39], who propose a novel method called *Faithfulness by Unlearning Reasoning Steps* (FUR). The central idea is that if a reasoning step is genuinely used by the model to derive its answer, then removing the corresponding information from the model’s parameters should alter its prediction.

For a given LLM M and a test question, the model first generates a CoT explanation. The explanation is then decomposed into individual reasoning steps. For each step, the corresponding knowledge is selectively unlearned from the model parameters, producing a fine-tuned model M' that is discouraged from predicting the removed reasoning step in the given context. Subsequently, both the original model M and the modified model M' are asked the same question again, this time without requesting a CoT explanation, and their predictions are compared.

For the proposed ff-hard metric, which evaluates the faithfulness of the entire CoT, the result is binary. A score of 1 indicates that the models disagree on the most likely answer, suggesting high faithfulness, whereas a score of 0 indicates that the prediction remains unchanged and the explanation likely does not reflect the true reasoning process. One major strength of this approach is that it avoids several pitfalls associated with context perturbation methods, such as the model reconstructing the modified context. However, the authors themselves note that high recall cannot be guaranteed: if unlearning a reasoning step does not alter the prediction, this may indicate failed unlearning, the existence of multiple reasoning paths, or the possibility that the step was not actually used. Furthermore, the FUR approach is not applicable to closed-source API-only models.

2.3 Causal Concept Faithfulness Metric

2.3.1 Introduction

Building upon the idea of counterfactual edits pioneered by Atanasova et al. [4] and Siegel et al. [34], Matton et al. [24] proposed an updated evaluation metric in their 2025 paper titled *Walk the Talk? Measuring the Faithfulness of Large Language Model Explanations*. Unlike previous approaches, their method called the *Causal Concept Faithfulness Metric* operates on higher-level concepts such as *gender*, *race* or *time of day* instead of low-level tokens or words. Furthermore, the method not only produces a quantitative score but also exposes semantic patterns related to faithfulness.

The authors define faithfulness as the alignment between truly influential concepts in the model’s prediction and the importance assigned to those concepts by the model itself within its explanation. To estimate the causal influence of concepts, an auxiliary LLM either replaces or removes concept values from the original input. As in the CCT metric, the original and counterfactual questions are then presented to the LLM and prediction distributions are collected, yielding estimates of the true causal effect. Subsequently, another helper LLM extracts the concepts deemed important

within the explanation, enabling the calculation of the explanation-implied effect. The Pearson correlation between the true causal effects and the explanation-implied effects then serves as the final faithfulness score.

One limitation of this approach lies in the LLM-based automation of counterfactual generation and explanation extraction, which may introduce errors and risks of circularity within the evaluation process. Furthermore, the metric does not perform well when concepts are strongly correlated. Since this metric is central to the present thesis, it will be discussed in greater technical detail in Chapter 2.3.

2.3.2 Mathematical Aspects

Problem Setting

The goal is to evaluate the faithfulness of explanations generated by an opaque large language model under evaluation, denoted by M , for a dataset of questions $X = \{x_1, \dots, x_N\}$. Matton et al. formulate their metric for so-called *context-based questions*, where each response r produced by the model consists of both a selected answer $y \in Y$ from a predefined set of answer options and a natural language explanation e , i.e.,

$$r = (y, e).$$

The underlying assumption is that the explanation implicitly or explicitly indicates which parts of the input influenced the model’s prediction. Rather than considering these parts as individual tokens or words, Matton et al. model them as high-level *concepts*. Consequently, the proposed metric measures the discrepancy between “the set of concepts that LLM explanations imply are influential and the set that truly are” [24].

Concepts and Categories

For each question x , a set of concepts

$$C = \{C_1, \dots, C_M\}$$

is identified. Unlike token-based approaches, concepts correspond to semantically meaningful properties of the input rather than low-level textual features. Each concept C_m can assume different possible values $c_m \in \mathbb{C}_m$, where $\mathbb{C}_m$ denotes the domain of all plausible values for concept C_m .

A key assumption of the metric—which will later be challenged in this thesis—is that concepts are “disentangled”. In other words, each concept can be modified independently without affecting the remaining concepts. This assumption allows causal interventions to be applied to individual concepts while keeping all others fixed.

Furthermore, every concept is assigned to a higher-level category from the set

$$K = \{K_1, \dots, K_L\},$$

for example *gender*, *ethnicity*, or *medical symptoms*. Besides improving the interpretability of the resulting analysis, these categories are also required for the Bayesian hierarchical model employed by Matton et al. to estimate causal concept effects through partial pooling.

Causal Concept Effects (CE)

To estimate the actual influence of a concept on the model’s prediction, Matton et al. generate counterfactual questions. A counterfactual,

$$x_{c_m \rightarrow c'_m},$$

is obtained by replacing the original value c_m of concept C_m with an alternative value c'_m , while leaving the remainder of the input unchanged. Let

$$\mathbb{C}'_m = \mathbb{C}_m \setminus \{c_m\}$$

denote the set of all alternative values of concept C_m .

The causal concept effect is then defined as the average Kullback–Leibler divergence between the model’s prediction distribution for the original question and each corresponding counterfactual:

$$CE(x, C_m) = \frac{1}{|\mathbb{C}'_m|} \sum_{c'_m \in \mathbb{C}'_m} D_{\text{KL}} \left(\mathbb{P}_M(Y \mid x_{c_m \rightarrow c'_m}) \parallel \mathbb{P}_M(Y \mid x) \right). \quad (2.1)$$

Intuitively, concepts whose modification substantially changes the model’s output receive a high causal concept effect, whereas concepts with little influence receive values close to zero.

Explanation-Implied Effects (EE)

While the causal concept effect measures the actual influence of a concept on the model’s prediction, the explanation-implied effect estimates how influential the model itself claims the concept to be within its explanation. If the explanations are faithful, concepts with large causal effects should be mentioned more frequently than concepts with only a negligible influence.

Matton et al. therefore define the explanation-implied effect as the probability that the model’s explanations for both the original question and its counterfactuals identify concept C_m as being causally relevant:

$$EE(x, C_m) = \frac{1}{|\mathbb{C}_m|} \sum_{c'_m \in \mathbb{C}_m} \mathbb{P}_M(C_m \in E \mid x_{c_m \rightarrow c'_m}). \quad (2.2)$$

In practice, this quantity is estimated by automatically extracting the concepts referenced in the generated explanations using an auxiliary LLM.

2.3.3 Causal Concept Faithfulness

After computing both the causal concept effects and the explanation-implied effects for every concept of a question, the final faithfulness score can be determined. Specifically, the Pearson correlation coefficient between the vectors of causal concept effects and explanation-implied effects is calculated:

$$F(x) = \text{PCC}(CE(x, C), EE(x, C)). \quad (2.3)$$

This yields a result within $[-1, 1]$. A high positive correlation indicates that concepts which genuinely influence the model’s prediction are also consistently identified as influential within the explanation, whereas a low or even negative correlation indicates unfaithfulness.

Finally, the dataset-level faithfulness of model M is obtained by averaging the question-level faithfulness scores over all questions in the dataset:

$$F(X) = \frac{1}{|X|} \sum_{x \in X} F(x). \quad (2.4)$$

2.3.4 Evaluation Pipeline

Extracting Concepts with Values and Categories

With a given question x drawn from the dataset X the concept set C has to be determined of all concepts that are included in the question. Since manual annotation by human reviewers would be tedious and impractical for an automated evaluation, an auxiliary LLM A is tasked with this work. It should be noted that A does not necessarily have to be the same model as the LLM-under-test M . The prompt instructs A to first identify the concepts C for the given question, as well as determine their respective category from the dataset-wide set K (however a list of possible categories is not supplied to the auxiliary llm). After that, in a second prompt, the LLM is then instructed to extract the current value of each identified concept as well as a plausible alternative value which results in $\mathbb{C}_m$ for every $C_m \in C$. All prompts within this pipeline feature examples to facilitate the LLM sticking to the desired output format. The result is then parsed and saved by the evaluation program and it can be continued to proceed to generate actual counterfactuals.

Generating counterfactual questions

By identifying the concepts involved in a question with their respective values, the cornerstone has been set for creating counterfactual question $x_{c_m \rightarrow c'_m}$. More specifically, to do so, the concept C_m , its current value c_m and the alternative value c'_m are passed down and A is instructed to overhaul x to $x_{c_m \rightarrow c'_m}$ while keeping all other concepts unchanged. So in summary, the auxiliary LLM A is queued to rephrase x with a new concept value. (with the least possible change to the structure or words not skew results, but the question should still be coherent).

Besides creating counterfactuals by changing the value of a concept C_m it might be, depending on the experimental setup and dataset, instead desirable to instead remove and obfuscate a concepts from a question. For doing so A is instructed to make C_m undeterminable from question x .

Collecting model responses

In order to be able to calculate the Causal Concept Effects (CE), its required to have access to the probability distributions of M 's choices for both original and all counterfactual questions $x_{c_m \rightarrow c'_m} : c'_m \in \mathbb{C}'_m$. To reconstruct those probability distributions a certain number of samples is being collected for every version of x by repeatedly prompting the LLM-under-test M . By doing so, not only are the answer choices obtained but also the explanations collected, which shall be subject to further analysis as well.

Estimating Explanation-Implied Effects

In order to get the Explanation-Implied Effects its needed to analyze the explanations on for looking if M indicates that C_m has a causal effect and might shift its prediction with alternative values. For automation, this step is also done by the auxillary LLM A . However: the auxillary LLM is instructed that one pitfall here is to assume that all concpets which are mentioned are influential and important. Instead A in method of Matton et al. is instructed to look for the context those concepts are mentioned in the explanation and not for there mere presence. (This is very akwardly written, need to rework that).

After analyzing all samples for all versions (original and counterfactuals) of a question x a DataFrame can be compiled for all $C_m \in x$ with the respective probability (averaged across all analyzed samples) that C_m is explained to have a causal effect. By using those probabilities its no difficulty to calculate the explanation-implied effects for the entire question (see equation 2.2).

Estimation Causal Concept Effects

Given that the model responses where already collected, it would be trivial to calculate the Kullback-Leibler divergence between the distribution of responses pre- and post-intervention for a given concept. However, Matton et al. point out that this approach can be inaccurate if only a small number of responses have been collected. By assuming that "that similar concepts have a similar magnitude of effect on LLM M 's answers within the context of a dataset" [24] they instead introduce a Bayesian hierarchical model within the context of the entire dataset. This allows the evaluation pipeline to "draw" information on the magnitude of concept effects of similar concepts across questions ($x \in X$). This multinomial logistic regression-approach is too rich and complex to be discussed in mathematical detail within this thesis, especially since it has only limited pertinence to the problem that should be adresssed.

By having made the calculation of the KL-divergence between both answer distributions more stable, its now possible to proceed to calculate the causal concept effects across all counterfactual values (see 2.1) for a given concept and question.

2.3.5 Estimating Causal Concept Faithfulness

This is the last step within Matton et al.’s’ evaluation pipeline. The Causal Concept faithfulness has been defined as the pearson correlation coefficient between the Causal Concept Effects and the Explanation-Implied Effects. However just with the CE-vector calculation the authors see a similar problem: for most questions the number of concepts involved is small, leading to only a handful of datapoints, and will make the results unreliable.

To solve this obstacle, Matton et al. suggest again to employ a Bayesian hierarchical model on the dataset-level, which "draws" data on faithfulness from other questions. They assume that faithfulness scores for questions within a dataset are rather similar due to a high degree of similarity of content. To make this linearly regression possible, the X and Y datapoints are prepared by normalizing their standard deviation to the same value. After that, a global regression parameter of the expected PCC is defined, serving as a prior on question-level regression coefficients. After fitting the model the posteriors of the question-regression coefficients are used as the question-level faithfulness. Extracting the global-regression coefficient would yield the $F(x)$ -value of the dataset.

Legacy Stuff

As discussed before, Matton et al. metric of causal concept faithfulness is based on the idea of measuring the difference between “the set of concepts that LLM explanations imply are influential and the set that truly are” [24]. For understanding how this metric works in detail, the evaluation pipeline is presented in the following step by step while keeping the notation introduced by Matton et al.

There is a model M which is an opaque black-box and a set of questions $X = \{x_1, \dots, x_N\}$. We can only then access discrete samples of the probability distribution $\mathbb{P}_M(R | x)$ of M to a question x .

2.3.6 The Case for this Metric

2.4 The Correlated Concepts Problem

Bibliography

- [1] Chirag Agarwal, Sree Harsha Tanneru, and Himabindu Lakkaraju. *Faithfulness vs. Plausibility: On the (Un)Reliability of Explanations from Large Language Models*. Mar. 14, 2024. DOI: 10.48550/arXiv.2402.04614. arXiv: 2402.04614 [cs]. URL: <http://arxiv.org/abs/2402.04614> (visited on 04/26/2026). Pre-published.
- [2] Sajid Ali et al. “Explainable Artificial Intelligence (XAI): What We Know and What Is Left to Attain Trustworthy Artificial Intelligence.” In: *Information Fusion* 99 (Nov. 1, 2023), p. 101805. ISSN: 1566-2535. DOI: 10.1016/j.inffus.2023.101805. URL: <https://www.sciencedirect.com/science/article/pii/S1566253523001148> (visited on 04/27/2026).
- [3] Dean Allemang and Juan Sequeda. “Increasing the Accuracy of LLM Question-Answering Systems with Ontologies.” In: *The Semantic Web – ISWC 2024*. Ed. by Gianluca Demartini et al. Cham: Springer Nature Switzerland, 2025, pp. 324–339. ISBN: 978-3-031-77847-6. DOI: 10.1007/978-3-031-77847-6_18.
- [4] Pepa Atanasova et al. “Faithfulness Tests for Natural Language Explanations.” In: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*. ACL 2023. Ed. by Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki. Toronto, Canada: Association for Computational Linguistics, July 2023, pp. 283–294. DOI: 10.18653/v1/2023.acl-short.25. URL: <https://aclanthology.org/2023.acl-short.25/> (visited on 05/06/2026).
- [5] *Blueprint for an AI Bill of Rights*. Oct. 2022. URL: <https://www.govinfo.gov/content/pkg/GOVPUB-PREX23-PURL-gpo193638/pdf/GOVPUB-PREX23-PURL-gpo193638.pdf> (visited on 04/25/2026).
- [6] Yanda Chen et al. “Do Models Explain Themselves? Counterfactual Simulatability of Natural Language Explanations.” In: *Proceedings of the 41st International Conference on Machine Learning*. Vol. 235. ICML’24. Vienna, Austria: JMLR.org, July 21, 2024, pp. 7880–7904.
- [7] Michael Mayowa Farayola et al. “Fairness of AI in Predicting the Risk of Recidivism: Review and Phase Mapping of AI Fairness Techniques.” In: *Proceedings of the 18th International Conference on Availability, Reliability and Security*. ARES ’23. New York, NY, USA: Association for Computing Machinery, Aug. 29, 2023, pp. 1–10. ISBN: 979-8-4007-0772-8. DOI: 10.1145/3600160.3605033. URL: <https://doi.org/10.1145/3600160.3605033> (visited on 04/27/2026).

- [8] Moncef Garouani et al. “Investigating the Duality of Interpretability and Explainability in Machine Learning.” In: *2024 IEEE 36th International Conference on Tools with Artificial Intelligence (ICTAI)*. 2024 IEEE 36th International Conference on Tools with Artificial Intelligence (ICTAI). Herndon, VA, USA: IEEE, Oct. 28, 2024, pp. 861–867. ISBN: 979-8-3315-2723-5. DOI: 10.1109/ICTAI62512.2024.00125. URL: <https://ieeexplore.ieee.org/document/10849387/> (visited on 04/26/2026).
- [9] Francisco Herrera. *Making Sense of the Unsensible: Reflection, Survey, and Challenges for XAI in Large Language Models Toward Human-Centered AI*. Comment: 30 pages, 1 figure. May 18, 2025. DOI: 10.48550/arXiv.2505.20305. arXiv: 2505.20305 [cs]. URL: <http://arxiv.org/abs/2505.20305> (visited on 04/25/2026). Pre-published.
- [10] Francisco Herrera. “Reflections and Attentiveness on eXplainable Artificial Intelligence (XAI). The Journey Ahead from Criticisms to Human–AI Collaboration.” In: *Information Fusion* 121 (Sept. 1, 2025), p. 103133. ISSN: 1566-2535. DOI: 10.1016/j.inffus.2025.103133. URL: <https://www.sciencedirect.com/science/article/pii/S1566253525002064> (visited on 04/25/2026).
- [11] Robert R. Hoffman et al. “Explainable AI: Roles and Stakeholders, Desires and Challenges.” In: *Frontiers in Computer Science* 5 (Aug. 17, 2023). ISSN: 2624-9898. DOI: 10.3389/fcomp.2023.1117848. URL: <https://www.frontiersin.org/journals/computer-science/articles/10.3389/fcomp.2023.1117848/full> (visited on 04/25/2026).
- [12] Weiche Hsieh et al. *A Comprehensive Guide to Explainable AI: From Classical Models to LLMs*. Dec. 8, 2024. DOI: 10.48550/arXiv.2412.00800. arXiv: 2412.00800 [cs]. URL: <http://arxiv.org/abs/2412.00800> (visited on 04/26/2026). Pre-published.
- [13] Chia-Chien Hung et al. “Walking a Tightrope – Evaluating Large Language Models in High-Risk Domains.” In: *Proceedings of the 1st GenBench Workshop on (Benchmarking) Generalisation in NLP*. Ed. by Dieuwke Hupkes et al. Singapore: Association for Computational Linguistics, Dec. 2023, pp. 99–111. DOI: 10.18653/v1/2023.genbench-1.8. URL: <https://aclanthology.org/2023.genbench-1.8/> (visited on 04/26/2026).
- [14] *ISO/IEC TR 29119-11:2020*. ISO. URL: <https://www.iso.org/standard/79016.html> (visited on 04/26/2026).
- [15] Alon Jacovi and Yoav Goldberg. “Towards Faithfully Interpretable NLP Systems: How Should We Define and Evaluate Faithfulness?” In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. ACL 2020. Ed. by Dan Jurafsky et al. Online: Association for Computational Linguistics, July 2020, pp. 4198–4205. DOI: 10.18653/v1/2020.acl-main.386. URL: <https://aclanthology.org/2020.acl-main.386/> (visited on 04/26/2026).

- [16] Jinqi Lai et al. “Large Language Models in Law: A Survey.” In: *AI Open* 5 (Jan. 1, 2024), pp. 181–196. ISSN: 2666-6510. DOI: 10.1016/j.aiopen.2024.09.002. URL: <https://www.sciencedirect.com/science/article/pii/S2666651024000172> (visited on 04/26/2026).
- [17] Tamera Lanham et al. *Measuring Faithfulness in Chain-of-Thought Reasoning*. July 17, 2023. DOI: 10.48550/arXiv.2307.13702. arXiv: 2307.13702 [cs]. URL: <http://arxiv.org/abs/2307.13702> (visited on 05/07/2026). Pre-published.
- [18] Weixin Liang et al. “The Widespread Adoption of Large Language Model-Assisted Writing across Society.” In: *Patterns* 6.12 (Dec. 12, 2025), p. 101366. ISSN: 2666-3899. DOI: 10.1016/j.patter.2025.101366. URL: <https://www.sciencedirect.com/science/article/pii/S2666389925002144> (visited on 04/26/2026).
- [19] Pedro Lopes et al. “XAI Systems Evaluation: A Review of Human and Computer-Centred Methods.” In: *Applied Sciences* 12.19 (Jan. 2022), p. 9423. ISSN: 2076-3417. DOI: 10.3390/app12199423. URL: <https://www.mdpi.com/2076-3417/12/19/9423> (visited on 04/27/2026).
- [20] Scott M. Lundberg and Su-In Lee. “A Unified Approach to Interpreting Model Predictions.” In: *Proceedings of the 31st International Conference on Neural Information Processing Systems*. NIPS’17. Red Hook, NY, USA: Curran Associates Inc., Dec. 4, 2017, pp. 4768–4777. ISBN: 978-1-5108-6096-4. URL: <https://dl.acm.org/doi/10.5555/3295222.3295230> (visited on 04/29/2026).
- [21] Haoyan Luo and Lucia Specia. *From Understanding to Utilization: A Survey on Explainability for Large Language Models*. Feb. 22, 2024. DOI: 10.48550/arXiv.2401.12874. arXiv: 2401.12874 [cs]. URL: <http://arxiv.org/abs/2401.12874> (visited on 04/29/2026). Pre-published.
- [22] Andreas Madsen, Sarath Chandar, and Siva Reddy. “Are Self-Explanations from Large Language Models Faithful?” In: *Findings of the Association for Computational Linguistics: ACL 2024*. Findings 2024. Ed. by Lun-Wei Ku, Andre Martins, and Vivek Srikumar. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 295–337. DOI: 10.18653/v1/2024.findings-acl.19. URL: <https://aclanthology.org/2024.findings-acl.19/> (visited on 04/26/2026).
- [23] Subhankar Maity and Manob Jyoti Saikia. “Large Language Models in Healthcare and Medical Applications: A Review.” In: *Bioengineering* 12.6 (June 10, 2025), p. 631. ISSN: 2306-5354. DOI: 10.3390/bioengineering12060631. PMID: 40564447. URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC12189880/> (visited on 04/26/2026).
- [24] Katie Matton, Robert Ness, and John Gutttag. “Walk the Talk? Measuring the Faithfulness of Large Language Model Explanations.” In: *The Thirteenth International Conference on Learning Representations*. Oct. 4, 2024. URL: <https://openreview.net/forum?id=4ub9gpx9xw> (visited on 04/24/2026).

- [25] Philip Mavrepis et al. *XAI for All: Can Large Language Models Simplify Explainable AI?* Jan. 23, 2024. DOI: 10.48550/arXiv.2401.13110. arXiv: 2401.13110 [cs]. URL: <http://arxiv.org/abs/2401.13110> (visited on 04/25/2026). Pre-published.
- [26] Melkamu Abay Mersha et al. *A Unified Framework with Novel Metrics for Evaluating the Effectiveness of XAI Techniques in LLMs*. Comment: arXiv admin note: substantial text overlap with arXiv:2501.15374. Apr. 7, 2025. DOI: 10.48550/arXiv.2503.05050. arXiv: 2503.05050 [cs]. URL: <http://arxiv.org/abs/2503.05050> (visited on 04/26/2026). Pre-published.
- [27] Maja Nisevic, Arno Cuypers, and Jan De Bruyne. *Explainable AI: Can the AI Act and the GDPR Go out for a Date?* Jan. 15, 2024. DOI: 10.2139/ssrn.5056022. Social Science Research Network: 5056022. URL: <https://papers.ssrn.com/abstract=5056022> (visited on 04/26/2026). Pre-published.
- [28] Letitia Parcalabescu and Anette Frank. “On Measuring Faithfulness or Self-consistency of Natural Language Explanations.” In: *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). Bangkok, Thailand: Association for Computational Linguistics, 2024, pp. 6048–6089. DOI: 10.18653/v1/2024.acl-long.329. URL: <https://aclanthology.org/2024.acl-long.329> (visited on 05/06/2026).
- [29] P Jonathon Phillips et al. *Four Principles of Explainable Artificial Intelligence*. NIST IR 8312. Gaithersburg, MD: National Institute of Standards and Technology (U.S.), Sept. 29, 2021, NIST IR 8312. DOI: 10.6028/NIST.IR.8312. URL: <https://nvlpubs.nist.gov/nistpubs/ir/2021/NIST.IR.8312.pdf> (visited on 04/27/2026).
- [30] Tehreem Qamar and Narmeen Zakaria Bawany. “Understanding the Black-Box: Towards Interpretable and Reliable Deep Learning Models.” In: *PeerJ Computer Science* 9 (Nov. 29, 2023), e1629. ISSN: 2376-5992. DOI: 10.7717/peerj-cs.1629. URL: <https://peerj.com/articles/cs-1629> (visited on 04/26/2026).
- [31] *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying down Harmonised Rules on Artificial Intelligence and Amending Regulations (EC) No 300/2008, (EU) No 167/2013, (EU) No 168/2013, (EU) 2018/858, (EU) 2018/1139 and (EU) 2019/2144 and Directives 2014/90/EU, (EU) 2016/797 and (EU) 2020/1828 (Artificial Intelligence Act) (Text with EEA Relevance)*. June 13, 2024. URL: <http://data.europa.eu/eli/reg/2024/1689/oj> (visited on 04/25/2026).
- [32] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. “Why Should I Trust You?": Explaining the Predictions of Any Classifier.” In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. KDD '16: The 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. San Francisco California USA: ACM,

- Aug. 13, 2016, pp. 1135–1144. ISBN: 978-1-4503-4232-2. DOI: 10.1145/2939672.2939778. URL: <https://dl.acm.org/doi/10.1145/2939672.2939778> (visited on 04/27/2026).
- [33] Ribana Roscher et al. “Explainable Machine Learning for Scientific Insights and Discoveries.” In: *IEEE Access* 8 (2020), pp. 42200–42216. ISSN: 2169-3536. DOI: 10.1109/ACCESS.2020.2976199. URL: <https://ieeexplore.ieee.org/document/9007737> (visited on 04/26/2026).
 - [34] Noah Siegel et al. “The Probabilities Also Matter: A More Faithful Metric for Faithfulness of Free-Text Explanations in Large Language Models.” In: *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*. ACL 2024. Ed. by Lun-Wei Ku, Andre Martins, and Vivek Srikumar. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 530–546. DOI: 10.18653/v1/2024.acl-short.49. URL: <https://aclanthology.org/2024.acl-short.49/> (visited on 05/06/2026).
 - [35] Ilia Stepin et al. “A Survey of Contrastive and Counterfactual Explanation Generation Methods for Explainable Artificial Intelligence.” In: *IEEE Access* 9 (Jan. 1, 2021), pp. 11974–12001. DOI: 10.1109/ACCESS.2021.3051315.
 - [36] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. “Axiomatic Attribution for Deep Networks.” In: *Proceedings of the 34th International Conference on Machine Learning - Volume 70*. ICML’17. Sydney, NSW, Australia: JMLR.org, Aug. 6, 2017, pp. 3319–3328. URL: <https://dl.acm.org/doi/10.5555/3305890.3306024> (visited on 04/29/2026).
 - [37] Maxwell Szymanski, Vero Vanden Abeele, and Katrien Verbert. “Disentangling Stakeholder Role and Expertise in User-Centered Explainable AI.” In: *Proceedings of the 33rd ACM Conference on User Modeling, Adaptation and Personalization*. UMAP ’25. New York, NY, USA: Association for Computing Machinery, June 13, 2025, pp. 32–39. ISBN: 979-8-4007-1313-2. DOI: 10.1145/3699682.3728351. URL: <https://dl.acm.org/doi/10.1145/3699682.3728351> (visited on 04/25/2026).
 - [38] Miles Turpin et al. “Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting.” In: *Proceedings of the 37th International Conference on Neural Information Processing Systems*. NIPS ’23. Red Hook, NY, USA: Curran Associates Inc., Dec. 10, 2023, pp. 74952–74965.
 - [39] Martin Tutek et al. “Measuring Chain of Thought Faithfulness by Unlearning Reasoning Steps.” In: *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*. EMNLP 2025. Ed. by Christos Christodoulopoulos et al. Suzhou, China: Association for Computational Linguistics, Nov. 2025, pp. 9935–9960. ISBN: 979-8-89176-332-6. DOI: 10.18653/v1/2025.emnlp-main.504. URL: <https://aclanthology.org/2025.emnlp-main.504/> (visited on 04/29/2026).

- [40] Giulia Vilone and Luca Longo. “Notions of Explainability and Evaluation Approaches for Explainable Artificial Intelligence.” In: *Information Fusion* 76 (Dec. 1, 2021), pp. 89–106. ISSN: 1566-2535. DOI: 10.1016/j.inffus.2021.05.009. URL: <https://www.sciencedirect.com/science/article/pii/S1566253521001093> (visited on 04/25/2026).
- [41] Jason Wei et al. “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.” In: *Proceedings of the 36th International Conference on Neural Information Processing Systems*. NIPS ’22. Red Hook, NY, USA: Curran Associates Inc., Nov. 28, 2022, pp. 24824–24837. ISBN: 978-1-7138-7108-8.
- [42] Quanjun Zhang et al. “A Survey on Large Language Models for Software Engineering.” In: *Science China Information Sciences* 69.4 (Mar. 5, 2026), p. 141102. ISSN: 1869-1919. DOI: 10.1007/s11432-025-4670-0. URL: <https://doi.org/10.1007/s11432-025-4670-0> (visited on 04/26/2026).
- [43] Haiyan Zhao et al. “Explainability for Large Language Models: A Survey.” In: *ACM Trans. Intell. Syst. Technol.* 15.2 (Feb. 22, 2024), 20:1–20:38. ISSN: 2157-6904. DOI: 10.1145/3639372. URL: <https://dl.acm.org/doi/10.1145/3639372> (visited on 04/27/2026).
- [44] Zining Zhu et al. “Explanation in the Era of Large Language Models.” In: *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 5: Tutorial Abstracts)*. Ed. by Rui Zhang, Nathan Schneider, and Snigdha Chaturvedi. Mexico City, Mexico: Association for Computational Linguistics, June 2024, pp. 19–25. DOI: 10.18653/v1/2024.naacl-tutorials.3. URL: <https://aclanthology.org/2024.naacl-tutorials.3/> (visited on 04/26/2026).
- [45] Mimi Zou and Lu Zhang. “Navigating China’s Regulatory Approach to Generative Artificial Intelligence and Large Language Models.” In: *Cambridge Forum on AI: Law and Governance* 1 (Jan. 2025), e8. ISSN: 3033-3733. DOI: 10.1017/cfl.2024.4. URL: <https://www.cambridge.org/core/journals/cambridge-forum-on-ai-law-and-governance/article/navigating-chinas-regulatory-approach-to-generative-artificial-intelligence-and-large-language-models/969B2055997BF42DE693B7A1A1B4E8BA> (visited on 04/25/2026).